Installing and working with caffe – CPU Version for testing different nets.

Durvesh Pathak

Electrical & Computer Engineering

Submitted to Dr. Mohamed El - Sharkawy

Index

Abstract:

This report aims at accelerating the process of installing caffe framework to test different types of Neural Networks.

This Report outlines the basic process of installing caffe on Ubuntu 17.04 ubuntu. You can also find information on http://caffe.berkeleyvision.org/install_apt.html. We will install caffe and run mnist example to check the installation.

Please note that this is not the installation process for Training the network that has to be done using GPU's. I will shortly start compiling process on training squeezenet.

Caffe: caffe is a deep learning framework developed by Berkeley AR research team. Which we will be using to train and deploy our Squeezenet V1.0.

Pl NOTE : the operating system used is 17.04

Pre Requisites:

1. Git

Git is a crucial requirement to clone repo from Github following command will install latest git on your system.

$ sudo apt-get install git

2. Python 2.7

Most Ubuntu installation comes with python 2.7.x but if you want to work on virtual environments install anaconda and the procedure for installation is available on this link - > https://yangcha.github.io/Caffe-Conda/

3. Opencv 3.x

To install this package type In the following commands

$ pip install opencv-python

$ pip install pillow

$ sudo apt-get install python-tk

You will need the following dependencies to run caffe:

$ sudo apt-get install -y --no-install-recommends libboost-all-dev
$ sudo apt-get install libprotobuf-dev libleveldb-dev libsnappy-dev libopencv-dev libboost-all-dev
$ sudo apt-get install vim
$ sudo apt-get install libhdf5-serial-dev libgflags-dev libgoogle-glog-dev liblmdb-dev protobuf-compiler
$ sudo apt-get install libopenblas-dev
$ sudo apt-get install libatlas-dev-dev
$ sudo apt install python-pip
$ sudo pip install scikit-image protobuf

Next, we will have to install all the necessary Python packages, using *pip*. Navigate to *python* folder, and type the line below:

$ cd python(Paths may be different for you)
$ for req in $(cat requirements.txt); do sudo pip install $req; done

Installation:

The very first thing to do is install caffe packages.

Everything including caffe itself is packaged in 17.04 and higher versions. To install pre-compiled Caffe package, just do it by

$ sudo apt install caffe-cpu

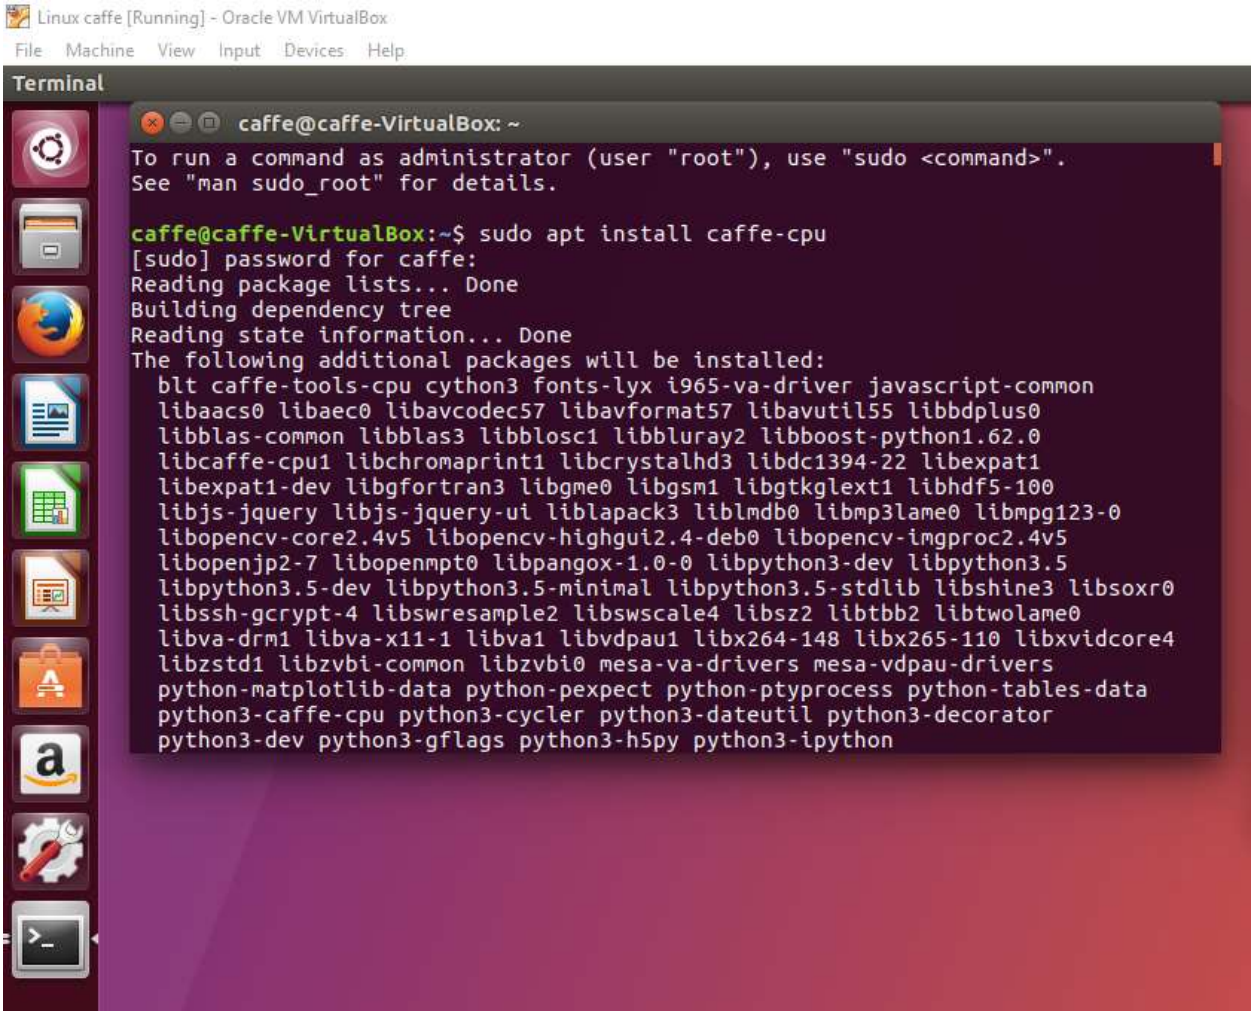

Fig: 1

This command will install the required caffe files on the local pc.

Next Clone the following repositories from GitHub.com I will be uploading Makefiles that will eliminate a lot of manual process later.

https://github.com/durveshpathak/SqueezeNet.git

https://github.com/durveshpathak/caffe.git (I will keep updating the python files asper our projects)

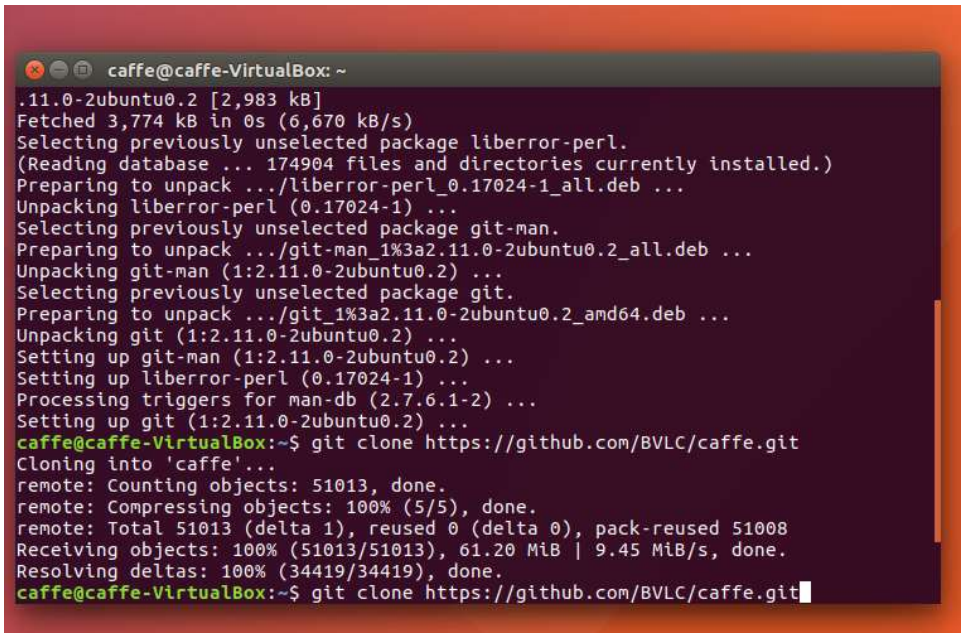

Fig: 2

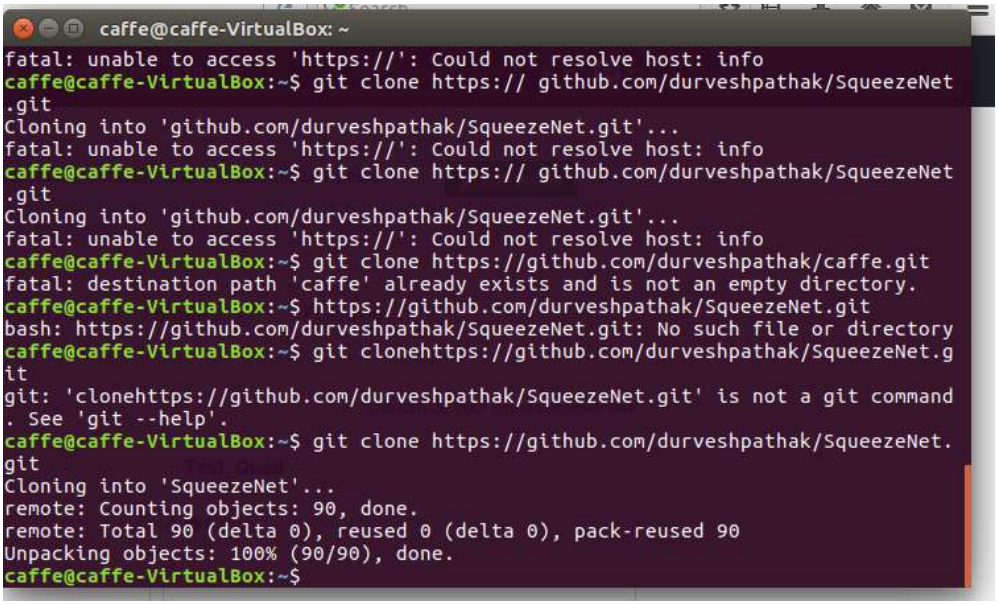

Fig: 3

Once you have cloned the above mentioned Repos from GitHub you should be able to see the following folders in your home directory home-→ caffe & Squeezenet

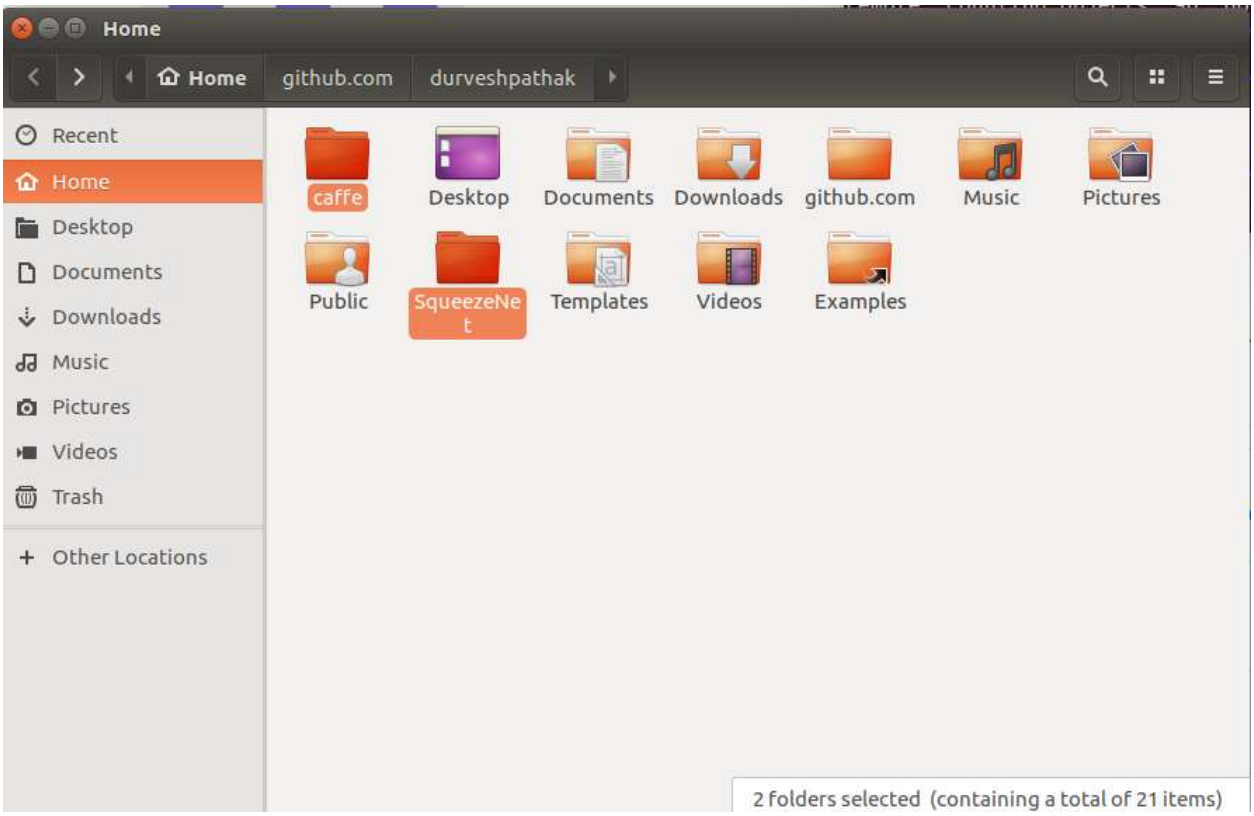

Fig: 4

Compilation:

Run the following commands in your terminal.

Change directory to caffe directory

$ cd caffe

$ ls (this will list all the directories & files.)

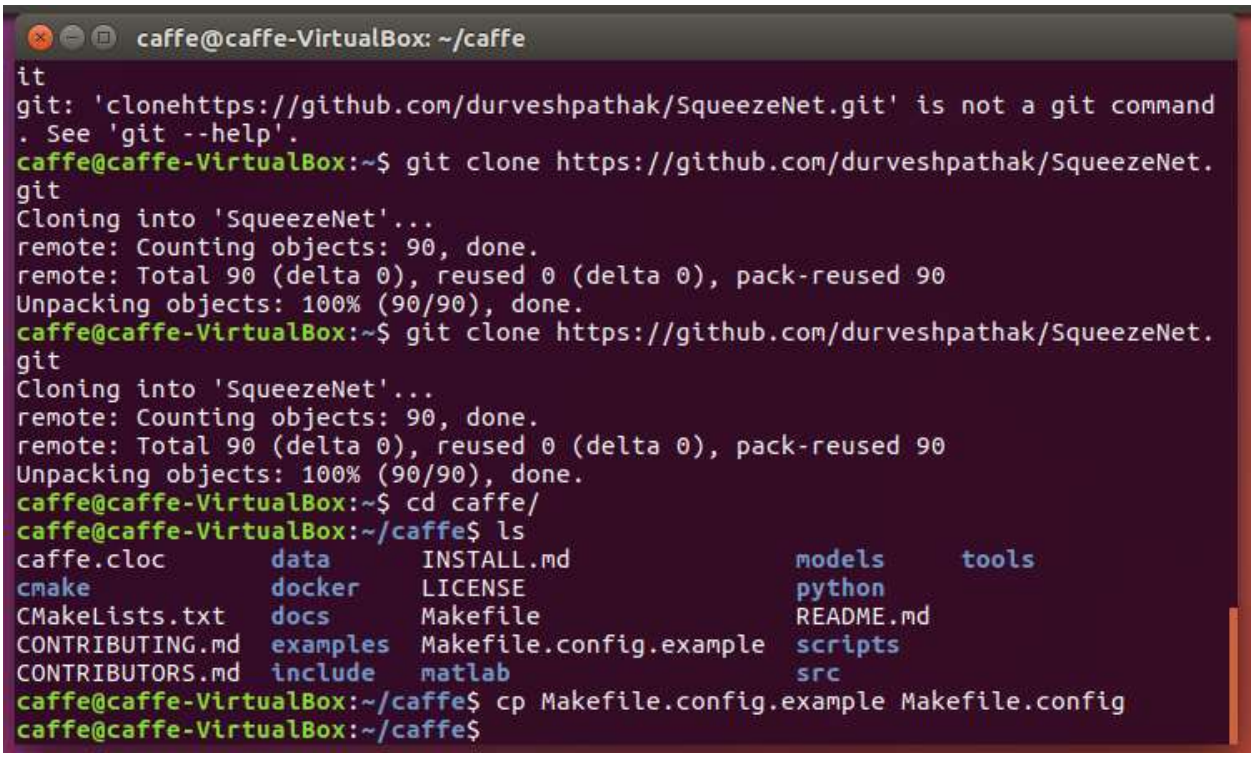

Fig: 5

Next we have to generate a MakeFile.config the default make file that comes with the repo cannot be directly used for CPU only modes.

The following command will make a copy of default MakeFile.config.example and save it as Makefile.config ← this is the file that is of utmost importance.

$ cp Makefile.config.example Makefile.config

After you create a makefile.config you might want to edit it based on whether you want to compile caffe for CPU or GPU. Since we are building caffe for CPU version we will change the following things in makefile.config & save the make file

Uncomment CPU_ONLY =1

Uncomment USE_LEVELDB =1

Uncomment USE_OPENCV =1

Uncomment USE_LMDB =1

Make sure when you are doing this you are in /caffe directory. The following command will start building caffe on you machine.

$ make all

You might receive an error!!!

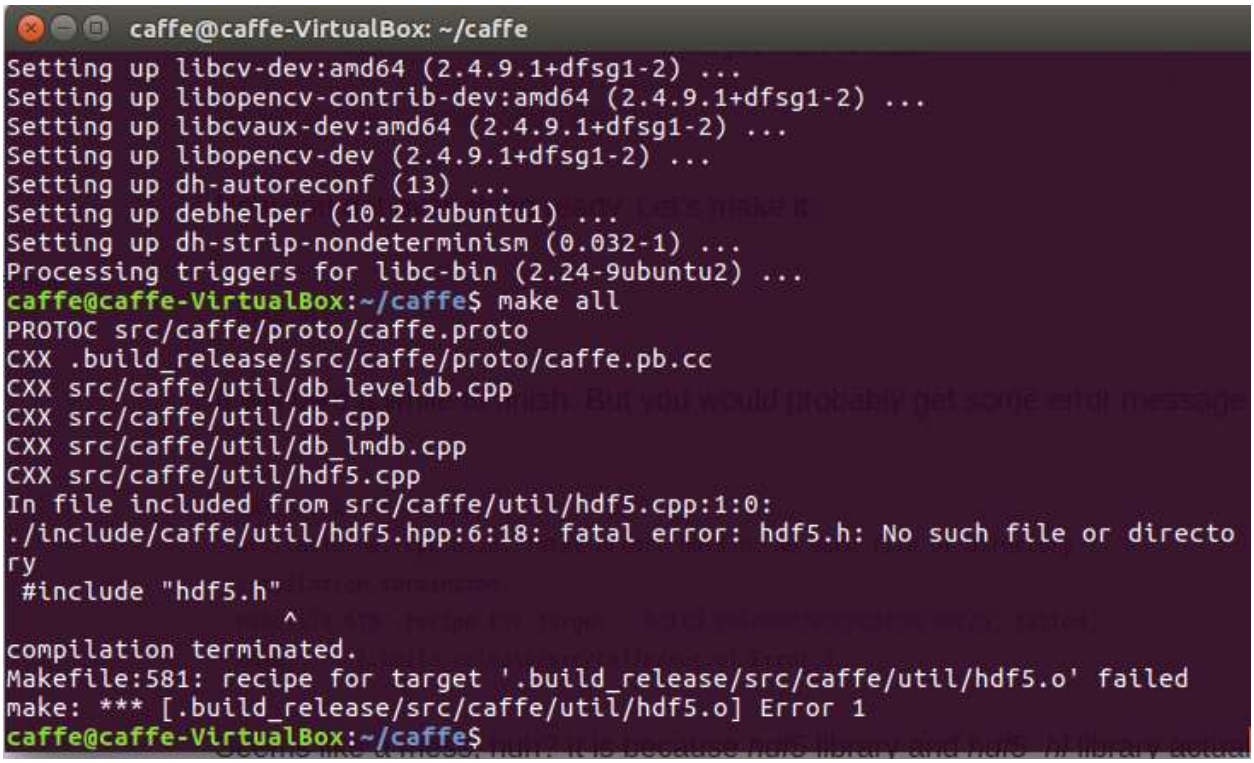

Fig: 6

```
CXX src/caffe/net.cpp
src/caffe/net.cpp:8:18: fatal error: hdf5.h: No such file or directory
compilation terminated.
Makefile:575: recipe for target '.build_release/src/caffe/net.o' failed
make: *** [.build_release/src/caffe/net.o] Error 1
```

Seems like a mess, huh? It is because hdf5 library and hdf5_hl library actually have a postfix serial in their names, the compiler cannot find them. To fix this, we just have to make a link to the actual files. Remember, we are not changing their names! But first, let's check out the actual name of the libraries. It may vary on your machines, though. To fix the issue you need to create a link file.

The following command will help you find the file.

$ cd /usr/lib/x86_64-linux-gnu/

$ ls -al

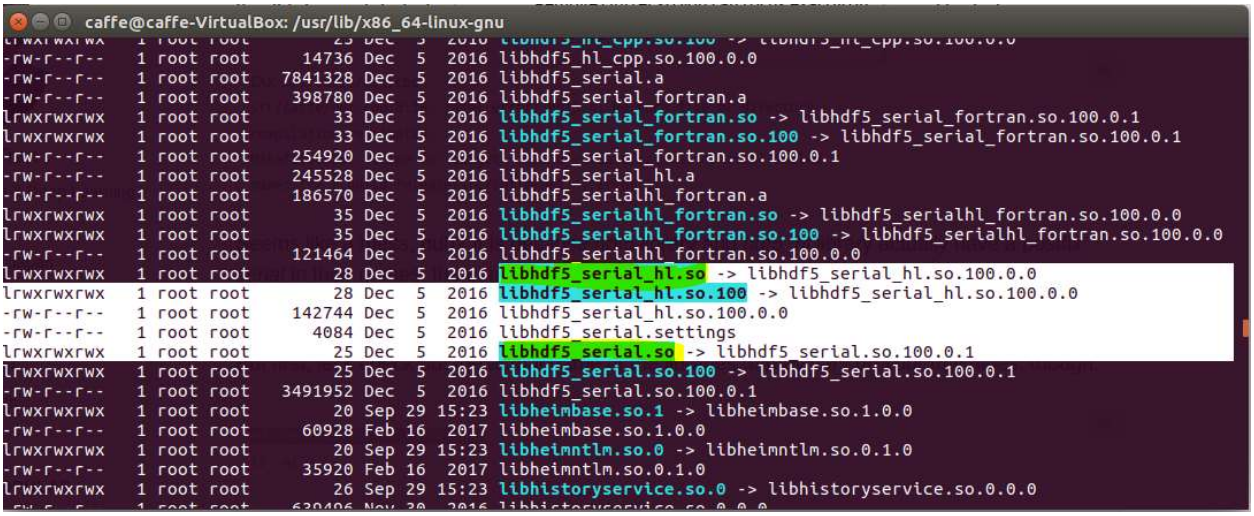

Fig: 7

You may find the two files like above. Note again that the version may be different. Just take note the ones you saw. Then we will make a link to them

$ sudo ln -s /usr/lib/x86_64-linux-gnu/libhdf5_serial.so. /usr/lib/x86_64-linux-gnu/libhdf5.so
$sudo ln -s /usr/lib/x86_64-linux-gnu/libhdf5_serial_hl.so.10.0.2 /usr/lib/x86_64-linux-gnu/libhdf5_hl.so

Please note: paths might differ a bit.

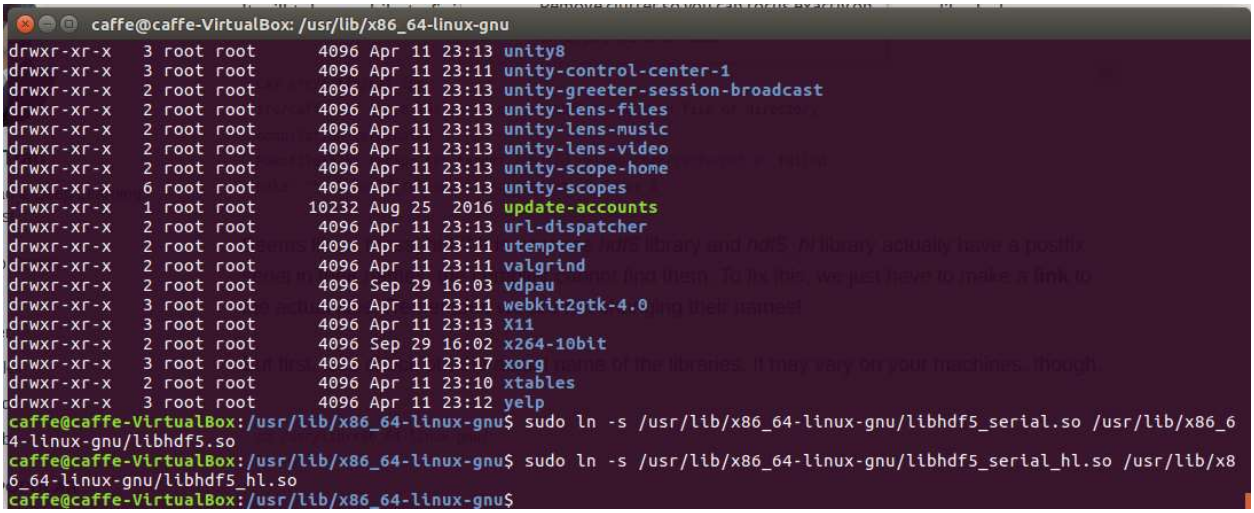

Fig: 8

After creating a link file type the following command

$ make all

If you get the following error, you need to make slight change in Makefile.config

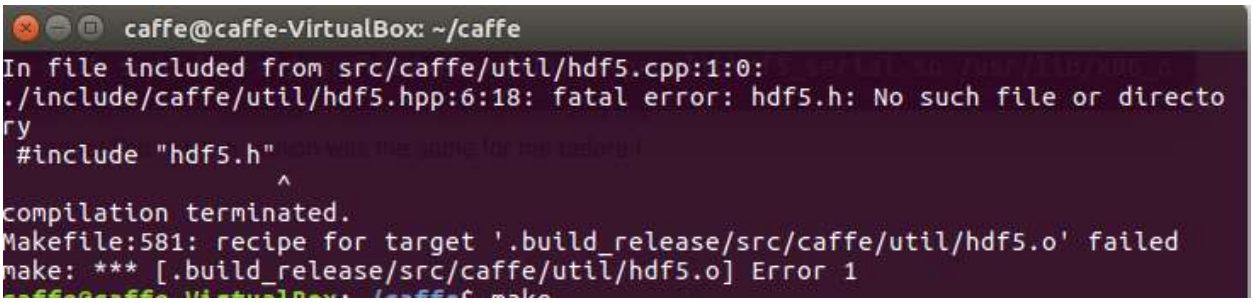

Fig:8

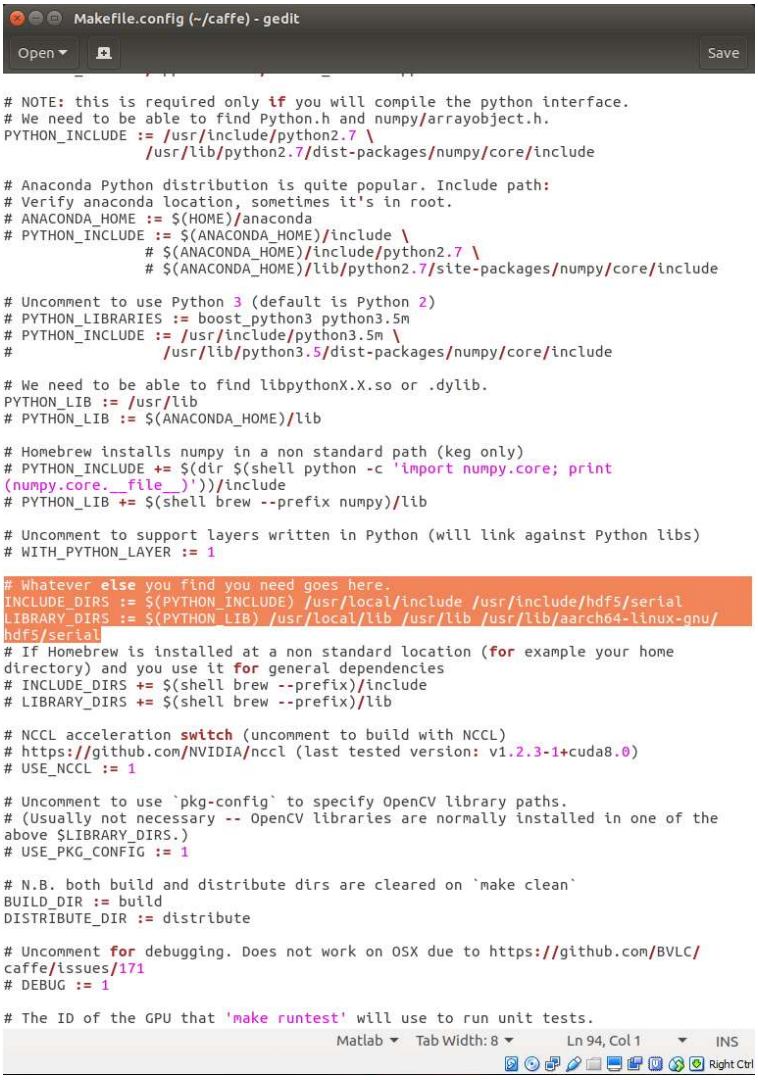

Fig: 9

Add the following line to makefile.config where it's highlighted

```
INCLUDE_DIRS := $(PYTHON_INCLUDE) /usr/local/include /usr/include/hdf5/serial/ and
LIBRARY_DIRS := $(PYTHON_LIB) /usr/local/lib /usr/lib /usr/lib/x86_64-linux-gnu/hdf5/serial/
```

Try running $ make all again. The caffe should compile.

Testing:

To test the caffe first type in the following commands

$ make test

And then

$ make runtest

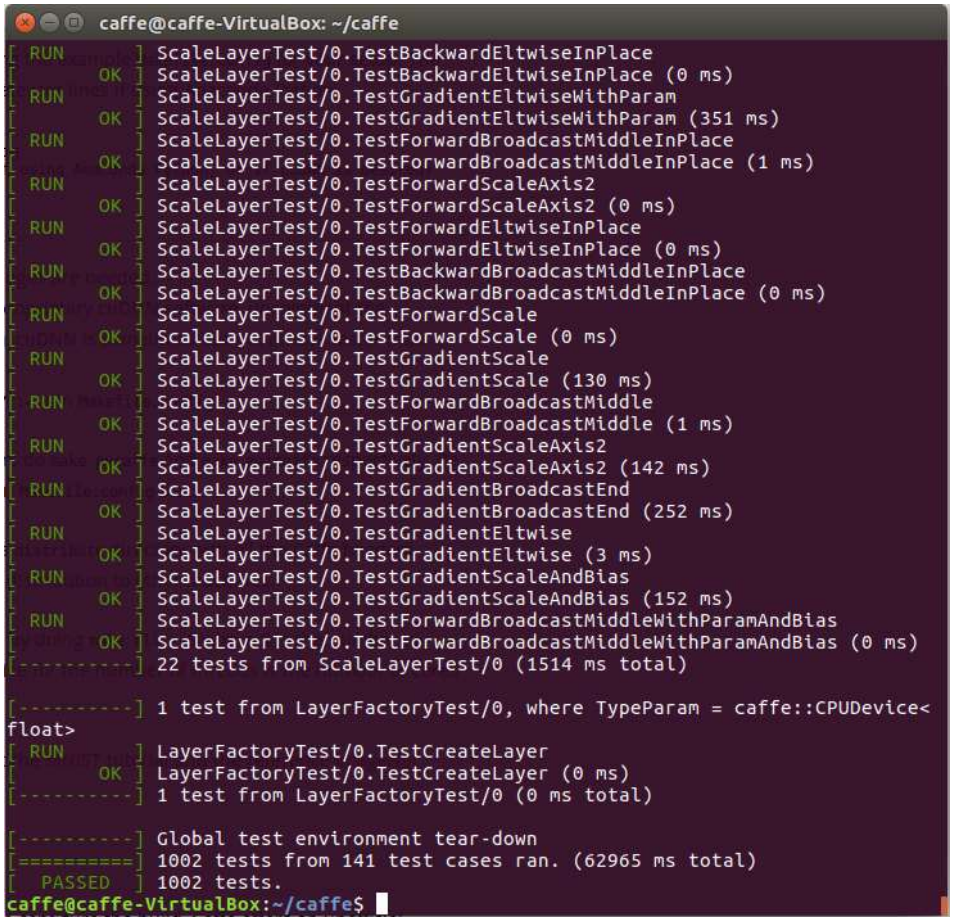

Fig: 10

If you saw something similar, then Congratulations! You have successfully installed Caffe! Now you can get your hands dirty with some real Deep Neural Network projects and become a part of Caffe community!

Next step is optional but I highly recommend because we are using Python for our works. We will compile the Python layer so that we can use *caffe* directly in our Python source code

Building Pycaffe:

Next step is to build pycaffe, for that type in the following command

$ make pycaffe

If you observe this error.

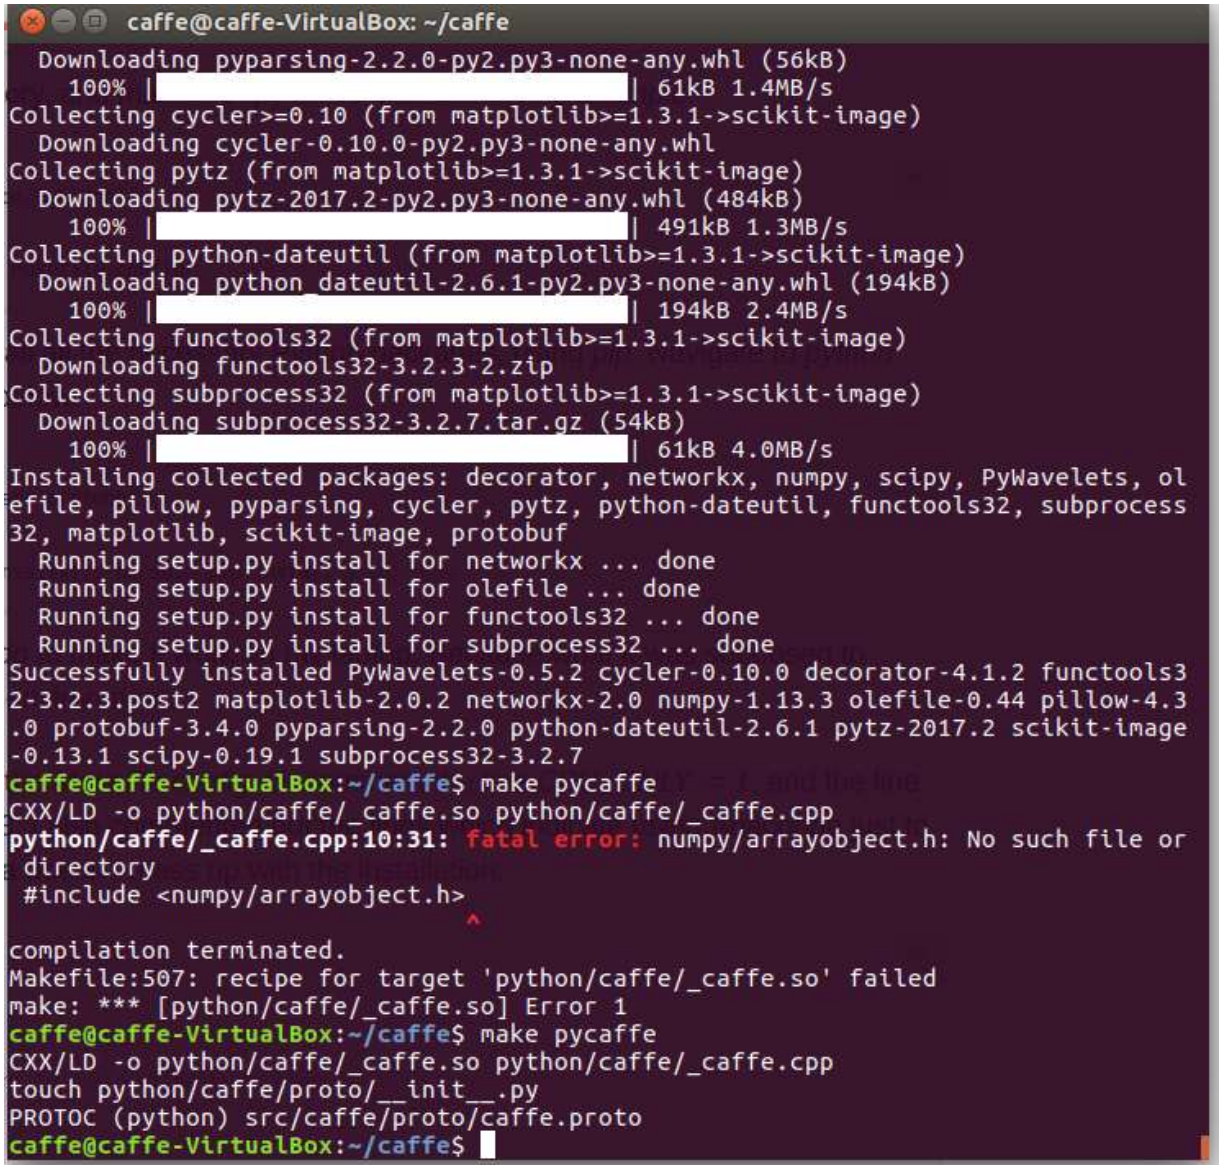

Fig: 11

```
CXX/LD -o python/caffe/_caffe.so python/caffe/_caffe.cpp
python/caffe/_caffe.cpp:10:31: fatal error: numpy/arrayobject.h: No such file or
directory
compilation terminated.
Makefile:501: recipe for target 'python/caffe/_caffe.so' failed
make: *** [python/caffe/_caffe.so] Error 1
```

make the following changes to makefile.config

```
PYTHON_INCLUDE := /usr/include/python2.7 \
/usr/lib/python2.7/dist-packages/numpy/core/include
```

to

```
PYTHON_INCLUDE := /usr/include/python2.7 \
/usr/local/lib/python2.7/dist-packages/numpy/core/include
```

Build pycaffe again!

To test the pycaffe move to python directory in caffe and start python environment and type import python this should not give you any error. You can also set this path to your environment variable.

$ cd caffe/python

$ python

>> import caffe

>>

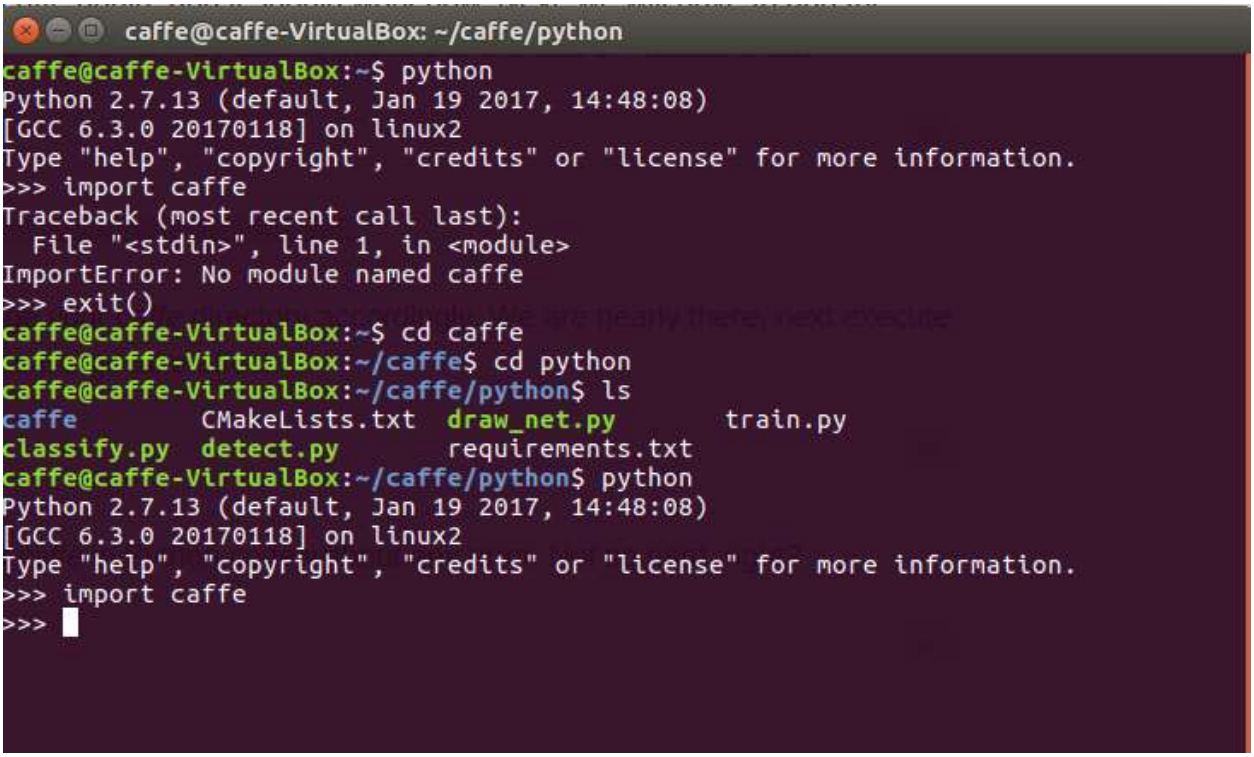

Fig: 12

Running examples:

But we are not done yet. Caffe provides us some examples of the most well-known models. We will use the LeNet model to train the MNIST dataset. Everything was already set up. All we have to do is just make it work:

$ cd /Downloads/caffe

$ ./data/mnist/get_mnist.sh

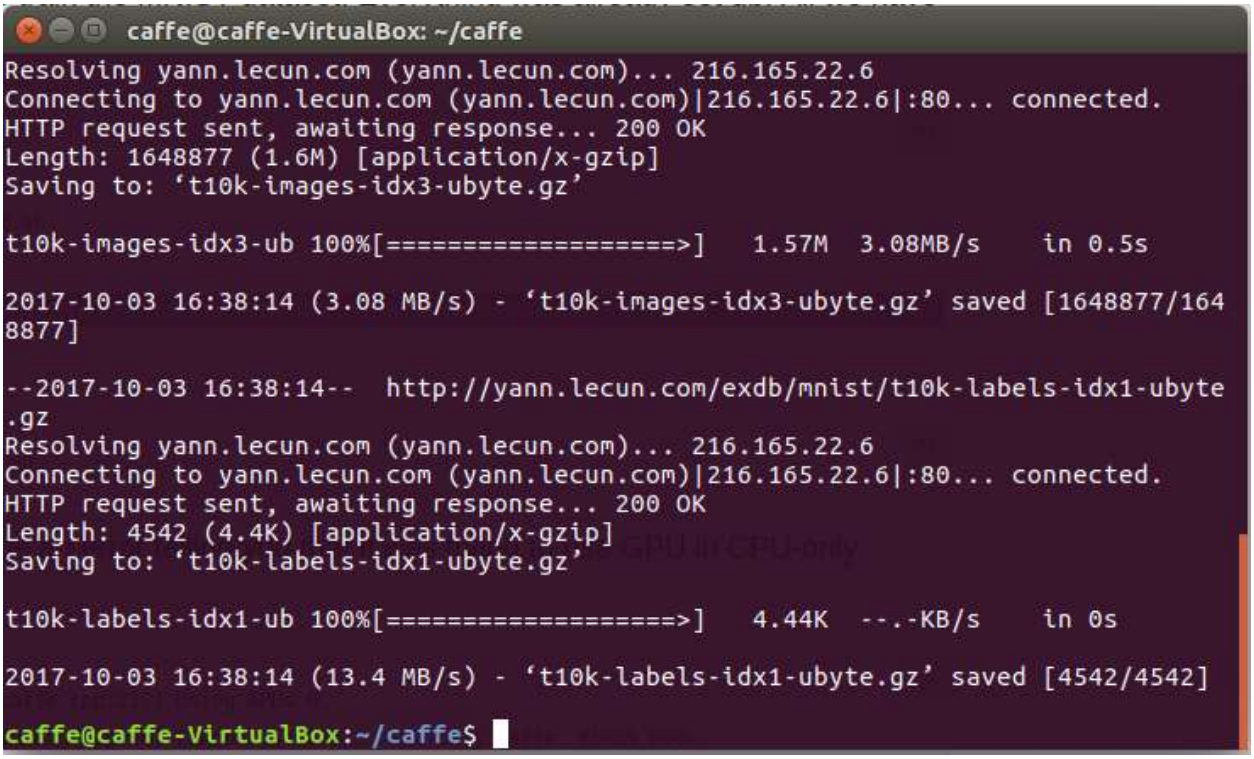

Fig: 13

$ ./examples/mnist/create_mnist.sh

$ ./examples/mnist/train_lenet.sh

If the above train command gives you the following error. Well, that's OK. We just have to apply a tiny fix to the file *examples/mnist/lenet_solver.prototxt*, replace *GPU* with **CPU**, and save it. Try to run the command above again, then everything should work just fine!

```
I1009 18:51:42.646926 22536 caffe.cpp:217] Using GPUs 0
F1009 18:51:42.647065 22536 common.cpp:66] Cannot use GPU in CPU-only Caffe: check
mode.
*** Check failure stack trace: ***
    @     0x7fd00383f5cd  google::LogMessage::Fail()
    @     0x7fd003841433  google::LogMessage::SendToLog()
    @     0x7fd00383f15b  google::LogMessage::Flush()
    @     0x7fd003841e1e  google::LogMessageFatal::~LogMessageFatal()
    @     0x7fd003c38c00  caffe::Caffe::SetDevice()
    @           0x40ad33  train()
    @           0x4071c0  main
    @     0x7fd0027b0830  __libc_start_main
    @           0x4079e9  _start
    @              (nil)  (unknown)
Aborted (core dumped)
```

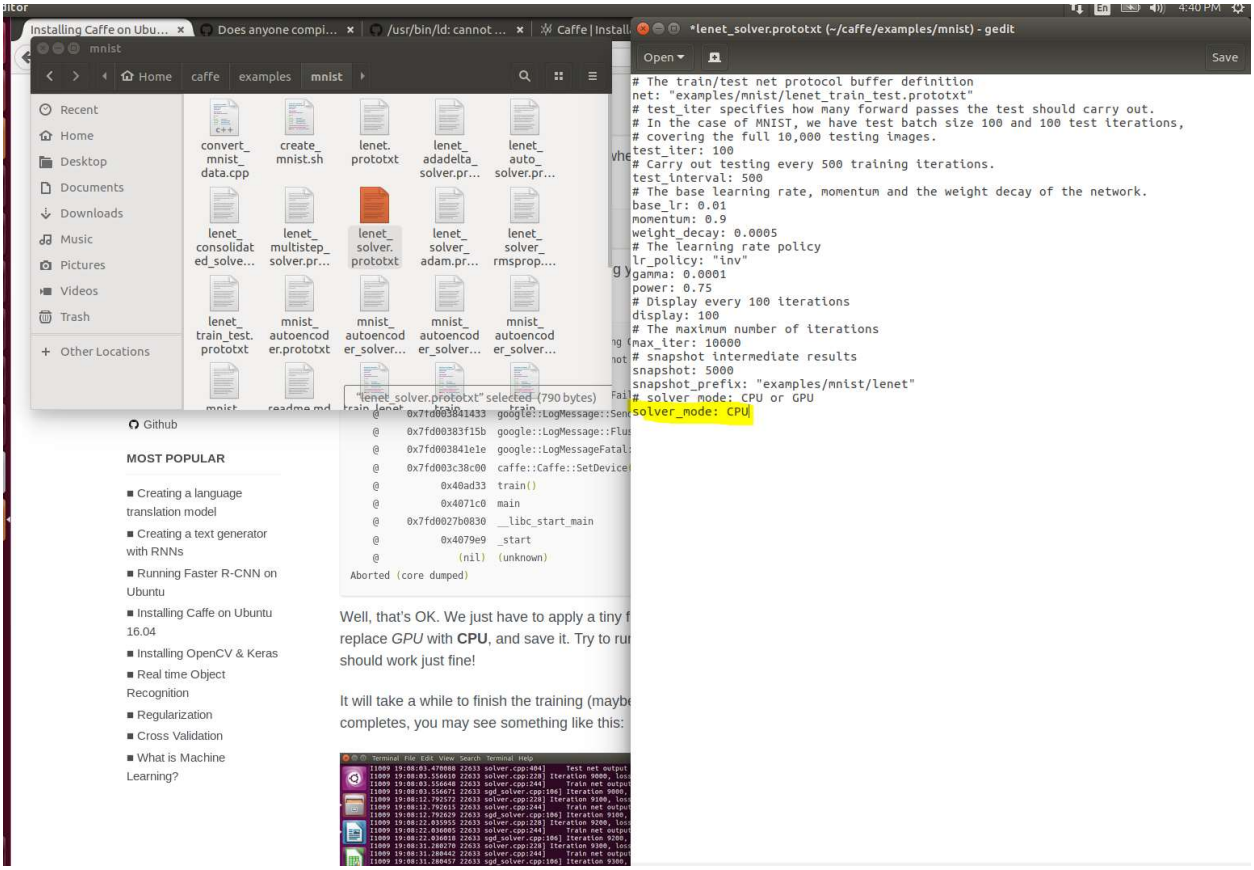

Fig: 14

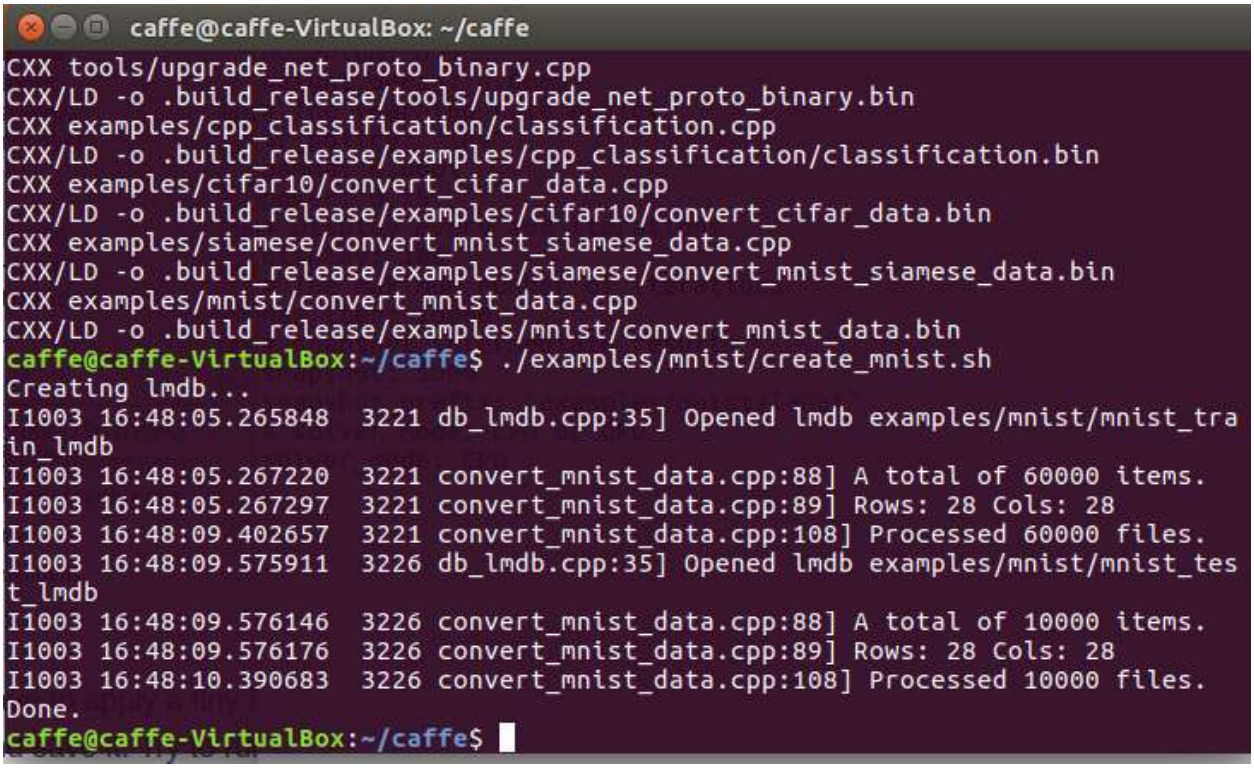

Fig: 15

It will take a while to finish the training (maybe long since we are using CPU). When it completes, you may see something like this:

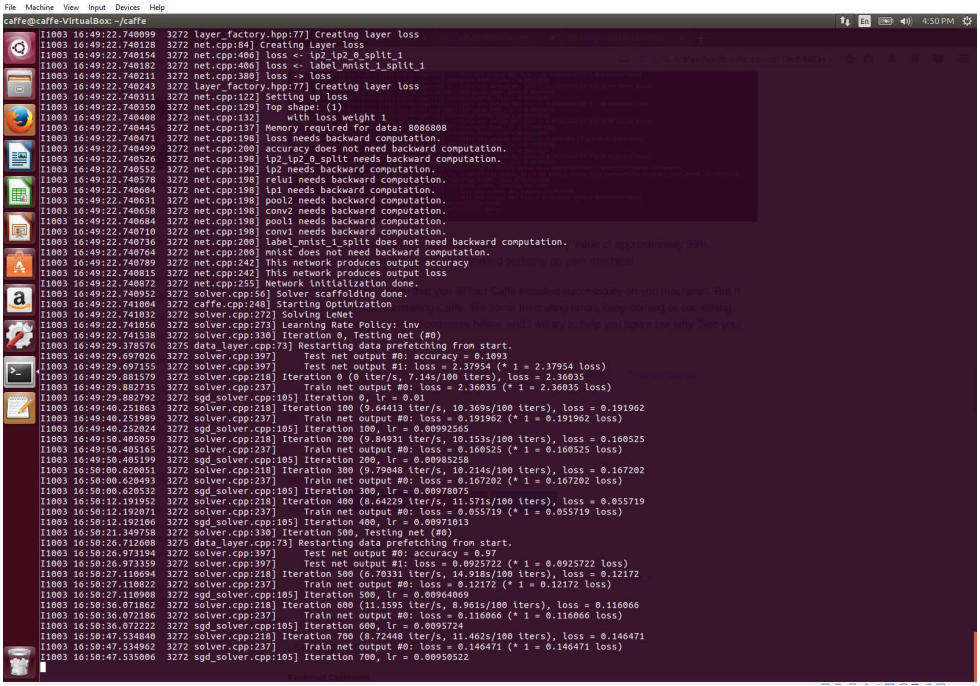

Fig: 16

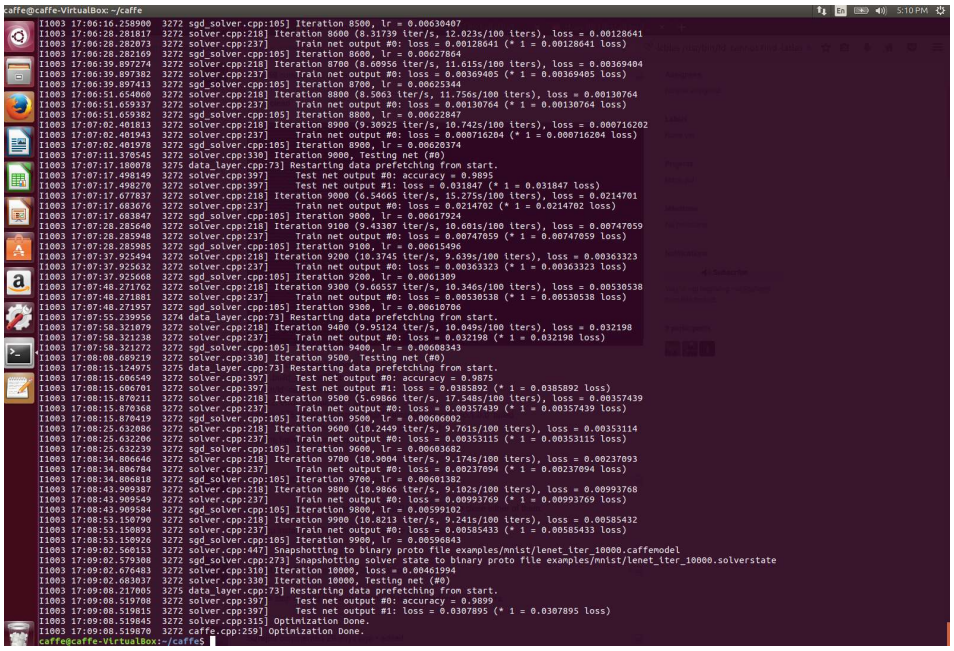

Fig: 17

Network is trained!!

So what have we trained!!!

I never explained what we trained, but since we have trained it let dig a little deeper to understand, what we have trained.

The simple answer is, we have trained Lenet network it's a type of Convolutional Neural Network that is popularly used for classification secondly we have trained the network with mnist data base which contains 60,000 image of digits 0,1,2….9.

So, since we have trained the network its time to run some example that we can visualize the output.

We will use a simple python script

Step 1: open a text editor.

Step 2: clone the following repo from github https://github.com/durveshpathak/caffe-mnist-test.git this will download a python code and image files.

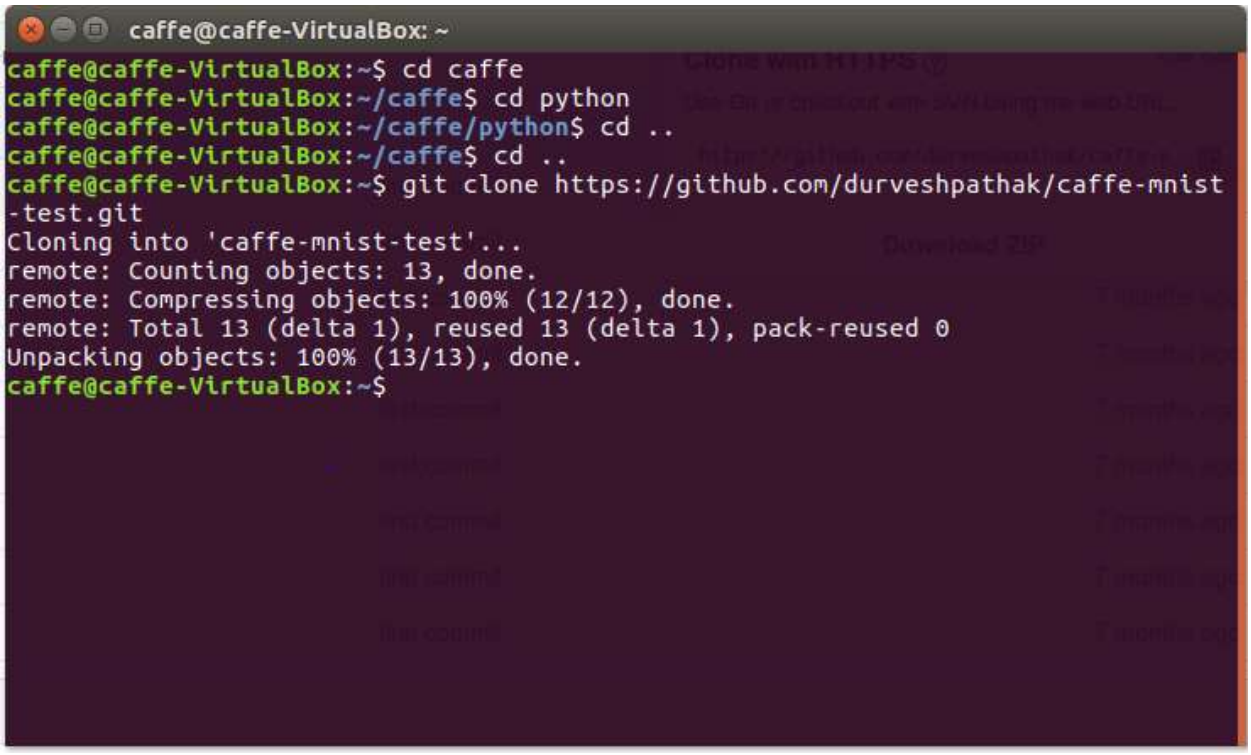

Step 3: Paste the files inside the downloaded folder in caffe/python directory

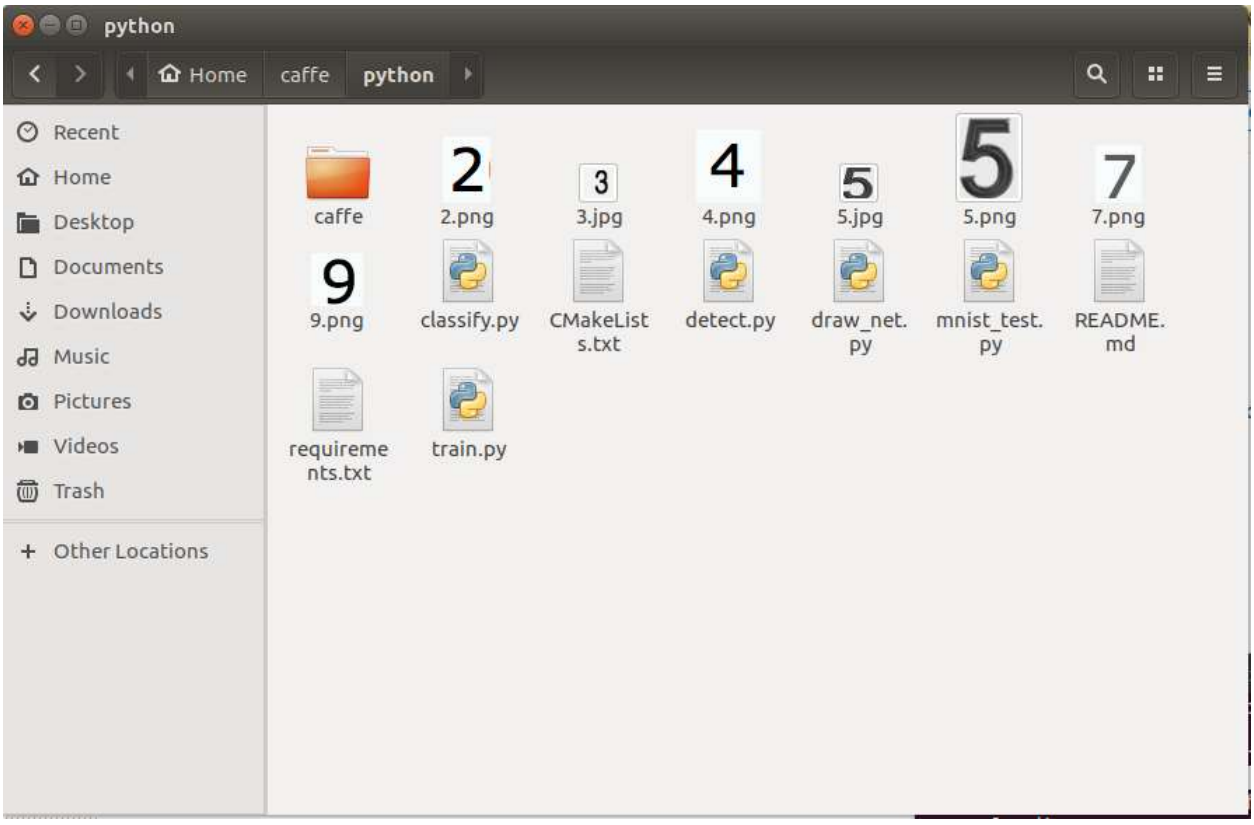

After pasting the files type in the following commands

$ python mnist_test.py

Conclusion : We can see that if we provide the png or jpg file for a specific digit the Lenet recognizes the digit and provides the prediction. Please make sure that you provide the right extension of the file. And Image and python file should be in same folder.

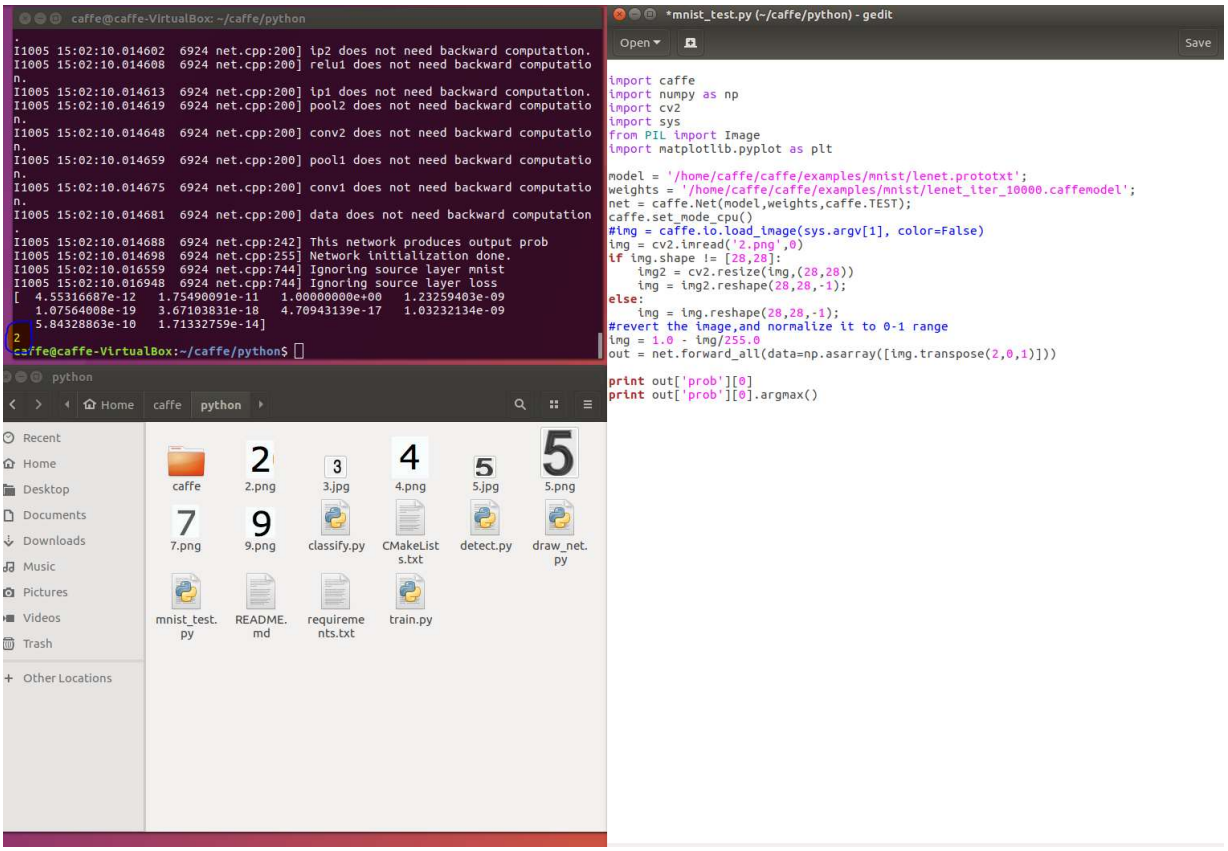
caffe@caffe-VirtualBox: ~/caffe/python
I1005 15:02:10.014602 6924 net.cpp:200] ip2 does not need backward computation.
I1005 15:02:10.014608 6924 net.cpp:200] relu1 does not need backward computation.
I1005 15:02:10.014613 6924 net.cpp:200] ip1 does not need backward computation.
I1005 15:02:10.014619 6924 net.cpp:200] pool2 does not need backward computation.
I1005 15:02:10.014648 6924 net.cpp:200] conv2 does not need backward computation.
I1005 15:02:10.014659 6924 net.cpp:200] pool1 does not need backward computation.
I1005 15:02:10.014675 6924 net.cpp:200] conv1 does not need backward computation.
I1005 15:02:10.014681 6924 net.cpp:200] data does not need backward computation.
I1005 15:02:10.014688 6924 net.cpp:242] This network produces output prob
I1005 15:02:10.014698 6924 net.cpp:255] Network initialization done.
I1005 15:02:10.016559 6924 net.cpp:744] Ignoring source layer mnist
I1005 15:02:10.016948 6924 net.cpp:744] Ignoring source layer loss
[4.55316687e-12 1.75490091e-11 1.00000000e+00 1.23259403e-09
1.07564008e-19 3.67103831e-18 4.70943139e-17 1.03232134e-09
5.84328863e-10 1.71332759e-14]
2
caffe@caffe-VirtualBox:~/caffe/python$
*mnist_test.py (~/caffe/python) - gedit
Open
Save
import caffe
import numpy as np
import cv2
import sys
from PIL import Image
import matplotlib.pyplot as plt
model = '/home/caffe/caffe/examples/mnist/lenet.prototxt';
weights = '/home/caffe/caffe/examples/mnist/lenet_iter_10000.caffemodel';
net = caffe.Net(model,weights,caffe.TEST);
caffe.set_mode_cpu()
#img = caffe.io.load_image(sys.argv[1], color=False)
img = cv2.imread('2.png',0)
if img.shape != [28,28]:
img2 = cv2.resize(img,(28,28))
img = img2.reshape(28,28,-1);
else:
img = img.reshape(28,28,-1);
#revert the image,and normalize it to 0-1 range
img = 1.0 - img/255.0
out = net.forward_all(data=np.asarray([img.transpose(2,0,1)]))
print out['prob'][0]
print out['prob'][0].argmax()
python
Home
caffe
python
Recent
Home
Desktop
Documents
Downloads
Music
Pictures
Videos
Trash
Other Locations
caffe
2.png
3.jpg
4.png
5.jpg
5.png
7.png
9.png
classify.py
CMakeLists.txt
detect.py
draw_net.py
mnist_test.py
README.md
requirements.txt
train.py

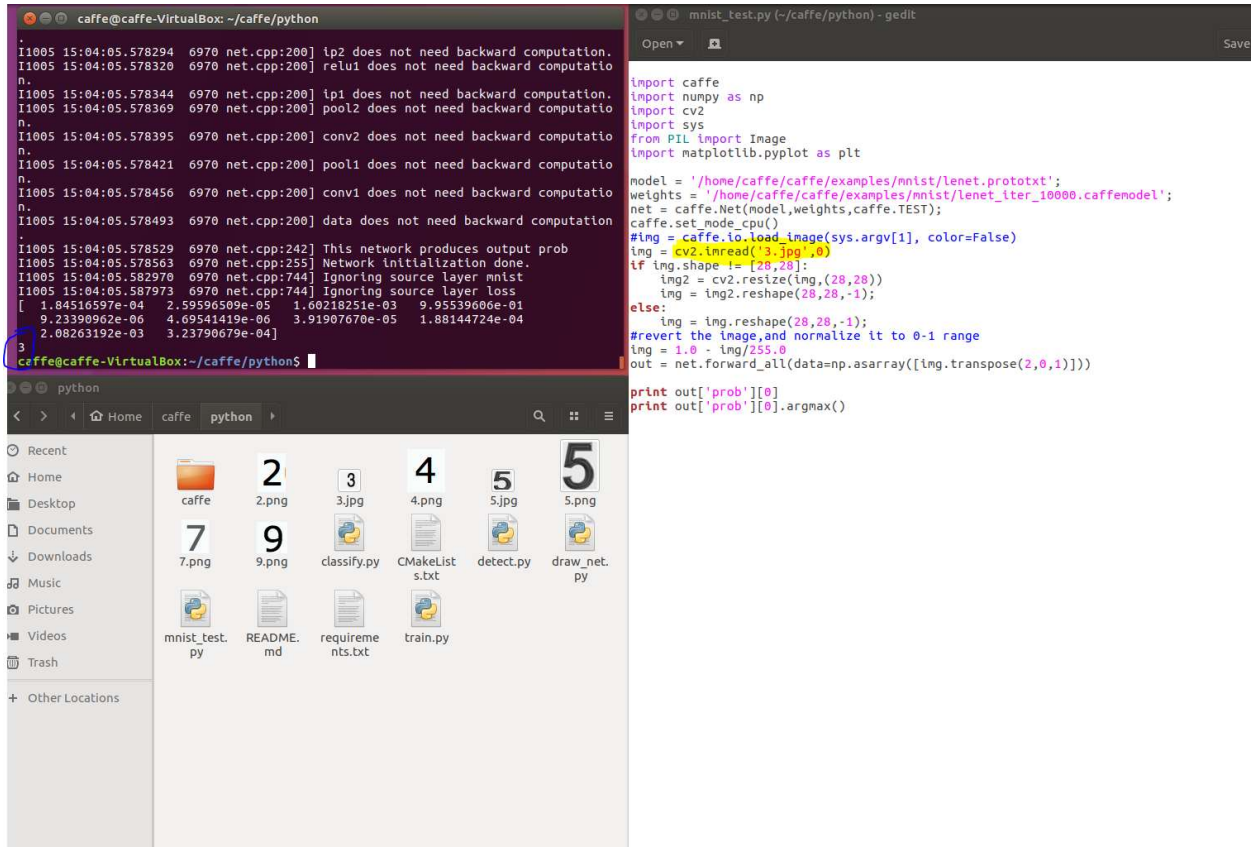
caffe@caffe-VirtualBox: ~/caffe/python
I1005 15:04:05.578294 6970 net.cpp:200] ip2 does not need backward computation.
I1005 15:04:05.578320 6970 net.cpp:200] relu1 does not need backward computation.
I1005 15:04:05.578344 6970 net.cpp:200] ip1 does not need backward computation.
I1005 15:04:05.578369 6970 net.cpp:200] pool2 does not need backward computation.
I1005 15:04:05.578395 6970 net.cpp:200] conv2 does not need backward computation.
I1005 15:04:05.578421 6970 net.cpp:200] pool1 does not need backward computation.
I1005 15:04:05.578456 6970 net.cpp:200] conv1 does not need backward computation.
I1005 15:04:05.578493 6970 net.cpp:200] data does not need backward computation.
I1005 15:04:05.578529 6970 net.cpp:242] This network produces output prob
I1005 15:04:05.578563 6970 net.cpp:255] Network initialization done.
I1005 15:04:05.582970 6970 net.cpp:744] Ignoring source layer mnist
I1005 15:04:05.587973 6970 net.cpp:744] Ignoring source layer loss
[1.84516597e-04 2.59596509e-05 1.60218251e-03 9.95539606e-01
9.23390962e-06 4.69541419e-06 3.91907670e-05 1.88144724e-04
2.08263192e-03 3.23790679e-04]
3
caffe@caffe-VirtualBox:~/caffe/python$
mnist_test.py (~/caffe/python) - gedit
Open
Save
import caffe
import numpy as np
import cv2
import sys
from PIL import Image
import matplotlib.pyplot as plt
model = '/home/caffe/caffe/examples/mnist/lenet.prototxt';
weights = '/home/caffe/caffe/examples/mnist/lenet_iter_10000.caffemodel';
net = caffe.Net(model,weights,caffe.TEST);
caffe.set_mode_cpu()
#img = caffe.io.load_image(sys.argv[1], color=False)
img = cv2.imread('3.jpg',0)
if img.shape != [28,28]:
img2 = cv2.resize(img,(28,28))
img = img2.reshape(28,28,-1);
else:
img = img.reshape(28,28,-1);
#revert the image,and normalize it to 0-1 range
img = 1.0 - img/255.0
out = net.forward_all(data=np.asarray([img.transpose(2,0,1)]))
print out['prob'][0]
print out['prob'][0].argmax()
python
Home
caffe
python
Recent
Home
Desktop
Documents
Downloads
Music
Pictures
Videos
Trash
Other Locations
caffe
2.png
3.jpg
4.png
5.jpg
5.png
7.png
9.png
classify.py
CMakeLists.txt
detect.py
draw_net.py
mnist_test.py
README.md
requirements.txt
train.py